

YAU, Chung Yiu

Email: cyau@umn.edu Webpage: <https://oscaryau525.github.io/>
Google Scholar: <https://scholar.google.com/citations?user=DXAnvT0AAAAJ>

RESEARCH EXPERIENCE

- ◇ (TMLR) (Arxiv) Data-parallel decentralized optimization. Neural network training on large-scale datasets in GPU clusters for parallel computation and minimal communication.
- ◇ (Arxiv) Accelerated optimization algorithms for deep learning.
- ◇ (ICML 2024) Optimization for contrastive learning and negative sampling.
- ◇ (ICML 2025) Large language model quantization for memory efficiency and inference speedup.

EDUCATION & ACADEMIA

Postdoctoral Associate Nov 2025 - Later
@ Department of Electrical and Computer Engineering, University of Minnesota.
Advisor: [Prof. Mingyi Hong](#).
◇ Research focus: optimization algorithms for large language model pre-training.

Ph.D. Systems Engineering & Engineering Management 2021 - 2025
@ The Chinese University of Hong Kong.
Supervisor: [Prof. Hoi-To Wai](#).
◇ Research focus: convergence analysis of new decentralized optimization algorithms with communication compression.

B.Sc. Computer Science 2017 - 2021
@ The Chinese University of Hong Kong.
◇ First Class Honour, ELITE Stream.

WORK EXPERIENCE

Applied Scientist Intern June - Sep 2024
@ Amazon Web Services, Santa Clara, California.
◇ Research on large language model quantization.
◇ Proposed an adaptive rotation method with [paper](#) presented at ICML 2025. Our algorithm fine-tunes a quantized LLM by quantization-aware training on Llama, Pythia and Qwen models.

Applied Scientist Intern June - August 2023
@ Amazon Web Services, Shanghai.
◇ Study the large batch dependency in contrastive learning.
◇ Proposed a small batch sampling algorithm with [paper](#) presented at ICML 2024. Our algorithm improves negative pair sampling and achieves exact convergence.

ACADEMIC ACTIVITIES

Invited Talk

Communication Efficient Decentralized Optimization

Sep 2024

@ Department of Electrical and Computer Engineering, University of Minnesota.

LLM Quantization-Aware Training

July 2025

@ INFORMS International 2025, Singapore.

Teaching Assistant

2021 - 2025

@ Department of Systems Engineering and Engineering Management, The Chinese University of Hong Kong.

- ◇ ENGG2440 Discrete Mathematics for Engineers
- ◇ FTEC2101 Optimization Methods

Voluntary Reviewer

@ conferences and journals:

- ◇ NeurIPS (2025 Top Reviewer, 2026), ICML (2024, 2025, 2026), ICLR (2025, 2026), TMLR
- ◇ IEEE TSP, ICASSP, TAC, L-CSS

.....

RESEARCH PUBLICATION

- [1] Chung-Yiu Yau, Dawei Li, Athanasios Glentis, Valentyn Boreiko, Hoi-To Wai, and Mingyi Hong. Ema-nesterov: Stabilizing nesterov’s lookahead for accelerated deep learning optimization. *arXiv preprint arXiv:2605.25395*, 2026.
- [2] Athanasios Glentis, Dawei Li, Chung-Yiu Yau, and Mingyi Hong. Revisiting the adam-sgd gap in llm pre-training: The role of large effective learning rates. *arXiv preprint arXiv:2605.17787*, 2026.
- [3] Chung-Yiu Yau, Haoming Liu, and Hoi-To Wai. A stochastic approximation approach for efficient decentralized optimization on random networks. *arXiv preprint arXiv:2410.18774v2*, 2025.
- [4] Haoming Liu, Chung-Yiu Yau, and Hoi-To Wai. Decentralized stochastic optimization over unreliable networks via two-timescales updates. *IEEE Transactions on Signal Processing*, 74:121–136, 2025.
- [5] Quan Wei*, Chung-Yiu Yau*, Hoi-To Wai, Dongyeop Kang, Youngsuk Park, Mingyi Hong, et al. Roste: An efficient quantization-aware supervised fine-tuning approach for large language models. *Proceedings of the 42st International Conference on Machine Learning*, 2025. *Equal contribution.
- [6] Haoming Liu, Chung-Yiu Yau, and Hoi-To Wai. A two-timescale primal-dual algorithm for decentralized optimization with compression. In *2025 IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*, 2025.
- [7] Chung-Yiu Yau, Hoi-To Wai, Parameswaran Raman, Soumajyoti Sarkar, and Mingyi Hong. EMC²: Efficient MCMC negative sampling for contrastive learning with global convergence. In *Proceedings of the 41st International Conference on Machine Learning*, pages 56966–56981. PMLR, 2024.

- [8] Chung-Yiu Yau and Hoi-To Wai. Fully stochastic distributed convex optimization on time-varying graph with compression. In *2023 62nd IEEE Conference on Decision and Control (CDC)*, pages 145–150. IEEE, 2023.
- [9] Xiaolu Wang, Chung-Yiu Yau, and Hoi-To Wai. Network effects in performative prediction games. In *International Conference on Machine Learning*, pages 36514–36540. PMLR, 2023.
- [10] Chung-Yiu Yau and Hoi-To Wai. Docom: Compressed decentralized optimization with near-optimal sample complexity. *Transactions on Machine Learning Research*, 2023.
- [11] Bingqing Song, Ioannis Tsaknakis, Chung-Yiu Yau, Hoi-To Wai, and Mingyi Hong. Distributed Optimization for Overparameterized Problems: Achieving Optimal Dimension Independent Communication Complexity. *Advances in Neural Information Processing Systems*, 2022.
- [12] Qiang Li, Chung-Yiu Yau, and Hoi-To Wai. Multi-agent Performative Prediction with Greedy Deployment and Consensus Seeking Agents. *Advances in Neural Information Processing Systems*, 2022.
- [13] Chung-Yiu Yau, Haoli Bai, Irwin King, and Michael R Lyu. DAP-BERT: Differentiable Architecture Pruning of BERT. In *International Conference on Neural Information Processing*, pages 367–378. Springer, 2021.